

Shellcoding, an introduction

Daniele Bellavista

*University of Bologna, Scuola di Ingegneria ed Architettura
Cesena Security Network and Applications*

December 11, 2012



Outline

- 1 Warm-Up
 - Shellcoding
 - Memory
 - Assembly
 - Sys-Calls
- 2 Writing shellcodes
 - Using a syscall
 - Data references
 - Execute a program
 - Reverse shell
- 3 Shellcode Optimization
 - String payload
- 4 References



Warm-Up

Materiale iniziale

- ▶ La teoria di questa presentazione vale per sistemi Unix, FreeBSD e Windows.
- ▶ Il materiale è tuttavia preparato per sistemi Linux 64 bit con kernel ≥ 2.6 e processore $>$ Pentium 4.
- ▶ Per chi non ha 64 bit sono disponibili anche gli esempi a 32 bit.
- ▶ Per chi non ha installato un OS Linux o non vuole sporcare il suo, sono disponibili due Virtual Machine:

32 bit: *sorry non ho fatto in tempo...*

64 bit: *sorry non ho fatto in tempo...*



What is shellcoding?

Exploiting software vulnerabilities

- ▶ Vulnerabilities like *BOFs* permit the insertion and the execution of a custom payload.
- ▶ The historical function of the payload is to spawn a shell (hence the name *shell-code*).
- ▶ If the exploited program runs as a privileged user, the payload can assume the control of the system.

Prerequisites

Basic notion of:

- ▶ Assembly language.
- ▶ Memory structure.
- ▶ System Calls.

System Memory

Memory management

- ▶ Memory management is a vast topic, so let's discuss from an high-level viewpoint.
- ▶ Thanks to the *Virtual Memory Management*, each process sees a flat addressing space of 128 TiB (4 GiB in 32 bit processors).
- ▶ Users can create sections inside the memory with read, write and/or execute permissions.
- ▶ The basic operations for memory management are *push* and *pop*.
 - push** Insert a 64-bit value into the top of the stack (the bottom).
 - pop** Retrieve a 64-bit value from the top of the stack (the bottom).




Figure : Linux 64 bit memory management



Program Code

Text segment and instruction pointer


- ▶ The code is saved in a RO and executable memory called *Text segment*.
- ▶ The current instruction is pointed by the *Instruction Pointer*.
- ▶ The instruction pointer is a value stored into the CPU.



Program Stack

- ▶ The stack is where function context, called frame, resides.
- ▶ The *Base pointer* points to the begin of the frame.
- ▶ The *Stack pointer* points to the first available memory location.

```
void f() {  
    int xx = 12; int yy = 23;  
    int zz = 24;  
    int sum = xx + yy + zz;  
}
```



Assembly Language

Low, low level programming

- ▶ Assembly is the final stage of the programming process.
- ▶ Each instruction is directly converted into an number.
- ▶ Writing the number on the CPU Command BUS cause the instruction to be executed.
- ▶ You can access to a set of 64 bit registers: *rax*, *rbx*, *rcx*, *rdx*, *rsi*, *rdi*, *r8*, *r9*, *r10*, *r11*, *r12*, *r13*, *r14*, *r15*
- ▶ As instruction pointer, base pointer and stack pointer the CPU uses respectively *rip*, *rbp* and *rsp*.



Assembly Language: function call

The call/ret instructions

call **address** push the RIP into the stack.

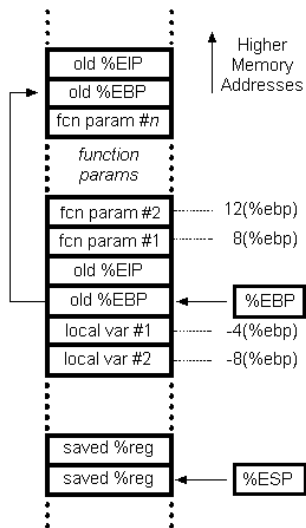
ret pop the RIP from the stack.

Creating a frame for the function

- ▶ The callee may need a stack for his own variables.
- ▶ *RBP* is saved (pushed) into the stack.
- ▶ The new *RBP* become the old *RSP*.
- ▶ Before doing *ret*, the opposite operations must be performed.



Assembly Language: function call and stack



Assembly Language: function call and parameters

Where should I put the parameters?

- ▶ If you are coding in pure ASM, you can do what you want.
- ▶ Else, you have to follow the conventions (see http://en.wikipedia.org/wiki/X86_calling_conventions#List_of_x86_calling_conventions).

Linux calling convention, parameter order

First parameters:

x86-64: rdi, rsi, rdx, rcx, r8, and r9. For system calls, r10 is used instead of rcx.

IA-32: ecx, edx. For system calls ebx, ecx, edx, esi, edi, ebp.

Other parameters are pushed into the stack.

System calls

Invoking the OS

- ▶ User's programs run in the exterior ring with *PL* 3, while kernel runs in the inner ring with *PL* 0: Kernel can access to the hardware (files, devices, ...), user not.
- ▶ Syscalls form the interface between user and kernel.
- ▶ E.g.: *open()*, *read()*, *write()*, *chmod()*, *exec()*, ...




Try to use Syscalls

Using assembly

- ▶ In order to ease the programmer duty, syscalls are identified by a number.
- ▶ In modern x86-64 processor, a syscall is invoked using the operation *syscall*, putting the id in the register *rax*.
- ▶ In modern x86 processor, a syscall is invoked using the operation *sysenter* but a lot of operation must be performed before, so in the following exercise we will use the old-fashion-way *int 80h*, which rises a software interrupt.

```
1  mov rdi , 0
2  mov rax , 60
3  syscall
```

```
1  mov ebx , 0
2  mov eax , 1
3  int 80h
```



Linux Syscalls

Linux

- ▶ Syscalls signatures: `/... kernel-sources... /include/linux/syscall.h`
- ▶ Syscalls numbers: `/usr/include/asm/unistd_64.h`

```
unistd_64.h
#define __NR_exit 60
syscall.h
asmlinkage long sys_exit(int error_code);
```



The final preparation

First steps...

- ▶ Testing platform: <http://dl.dropbox.com/u/16169203/data.zip>
- ▶ Use nasm to compile your assembly code
- ▶ Feed the tester with the output.

Exercise 0

- ▶ Create a shellcode that does nothing!
- ▶ Test it!



Using a syscall

Exercise 1: The exit syscall

- ▶ Use the exit syscall to terminate the program
- ▶ Change your code setting the return code to `13h`



Data reference I

The reference problem

- ▶ When you write a program you can use global data because, at compile time, a static address is associated.
- ▶ But your shellcode is not compiled with the program it's intended to run.
- ▶ You must use relative addressing, but before IA-64 it was not possible.



Data reference II

Old IA-32 way

- ▶ You use a trick: `jmp` just before the data location, then do a `call`. Et voilà! On the top of the stack there is the data address.

```
jmp message
run:
    pop ebx ; now ebx contains the string reference
    ; ... shellcode
message:
    call run
    db 'CeSeNA',0x0a,0
```



Data reference III

New IA-64 way

- ▶ IA-64 introduces the RIP relative addressing.

```
lea rdi, [rel message] ; now ebx contains  
                        ; the string reference  
; ... shellcode  
  
message db 'CeSeNA', 0x0a, 0
```



Data reference IV

Generic Way

- ▶ You can pop the string in hex format over the stack.
- ▶ The stack pointer is then the string reference.

```
push 0x000a414e ; 0x00, 0x0a, 'AN'  
push 0x65536543 ; 'eSeC'  
mov ebx, esp ; now ebx contains the string reference  
; ...
```



Data reference V

Exercise 2: print on stdout

- ▶ Get the exercise skeleton.
- ▶ Understand the code
- ▶ Write your own message



Execute a program


```
unistd_64.h
#define __NR_execve 59
syscall.h
int kernel_execve( const char *filename ,
                   const char *const argv[] ,
                   const char *const envp[] );
```

Exercise 3: A real shellcode, exec a shell!

4 HaXoRs

4 g33k

4 n00bz



Reverse shell I

One shellcode to rule them all

Step to execute:

- 1 Open a socket to the attacker server.
- 2 Duplicate the socket file descriptor into 0 and 1 (and optionally 2).
- 3 Exec a shell.

Step1: Create a socket

- ▶ The standard procedure involves socket creation and connection.
- ▶ `socket()` and `connect()` syscalls.
- ▶ The `sockaddr_in` structure requires port in network byte order (`htons()`) and ip address in numeric form (`inet_pton()`).
- ▶ The usual `sockaddr_in` size is 16.

Reverse shell II

```
long sys_socket(int domain,
                int type,
                int protocol);
long sys_connect(int fd,
                 struct sockaddr __user *,
                 int addrlen);
```

```
struct sockaddr_in {
short                sin_family; // AF_INET
unsigned short      sin_port; // network order (htons())
struct in_addr      sin_addr; // As 32 bit
char                sin_zero[8];
};
```



Reverse shell III

Step2: Duplicate the file descriptor

- ▶ The return value of `socket()` is the file descriptor.
- ▶ The syscall `dup2()` copy the file descriptor and close the destination.
- ▶ `dup2(fd, 0); dup2(fd, 1); dup2(fd, 2);`

```
long sys_dup2(unsigned int oldfd ,  
              unsigned int newfd)
```



Reverse shell IV

Step3: Exec a shell

Already seen ;)

Exercise 4: Reverse shelling

Remember to open a server listening into a known address! E.g.: `nc -l -p 1234`, preparing reverse shell for 127.0.0.1:1234

4 HaXoRs

4 g33k

4 n00bz



Shellcode optimization

Why optimization?

- ▶ When writing a shellcode, you must consider various factors. The most important are:
 - 1 How the vulnerable program receives the payload.
 - 2 How long the payload can be.
- ▶ The previous exercise solutions, for instance, are not suitable in most cases.
- ▶ The next slides will discuss about the *NULL* byte presence.
- ▶ About the payload length, it's all about the exploiter expertise.



String payload

zero-bytes problem

- ▶ If the payload is read as a string, C interprets a zero byte as string terminator, cutting the shellcode.
- ▶ Zero-bytes presence is caused by data and addresses:
 - ▶ `mov rax, 11h` is equivalent to `mov rax, 0000000000000011h`.
 - ▶ `lea rax, [rel message]` is equivalent to `lea rax, [rip + 0000...xxh]`.
 - ▶ `execve` for instance, requires a null terminated string and some null parameters.
- ▶ Solutions are quite straightforward:
 - ▶ Use `xor` operation to zero a register.
 - ▶ Use when possible smaller part of registers (e.g.: `rax` → `eax` → `ax` → `[ah,al]`)
 - ▶ Use `add` operation: immediate operators are not expanded.
 - ▶ Place not-null marker in strings and substitute them inside the code.
 - ▶ When using relative addressing, place the message above: offset will be negative [3].

Zero-byte removal example

```
; Set rax = 60h
xor rax, rax
mov al, 60
```

```
; Set to 0 a mem area
null db 'xxxx'
xor rbx, rbx
mov [rel null], ebx
```

```
; Set rdi = 12h
xor rdi, rdi
add rdi, 12h
```

```
; terminate string with 0
message db 'CeSeNA', 'x'
xor rbx, rbx
lea rdi, [rel message]
mov [rdi+7], bl
```

```
; Negative reference
message db 'CeSeNA', 'x'
lea rdi, [rel message]
```



References



Linux Foundation.
Linux documentation.
<http://linux.die.net/>.



Intel.
*Intel® 64 and IA-32 Architectures Software Developer's Manuale
Combined Volumes:1, 2A, 2B, 3C, 3A, 3B and 3C.*
Intel, August 2012.



Mark Loiseau.
64 bit linux shellcode.
<http://blog.markloiseau.com/2012/06/64-bit-linux-shellcode/>.



G. Adam Stanislav.
<http://www.int80h.org/>.

